

Scheduling and schedulers

Dr. Bystrov

School of Electrical Electronic and Computer Engineering
University of Newcastle upon Tyne

Introduction

- Remember concurrent programming? How is concurrency implemented in a system with a single processor?
- Remember concurrent models? How close to the “true concurrency” can we approach?
- An Operating System (OS) – is it incomprehensibly complex?
- Is it possible to guarantee compliance with the Real-Time constraints in some particular system?
 - What is Real-Time?
 - How to measure it?
 - Which scheduling algorithm to choose?
 - Is there the best scheduler for an application?

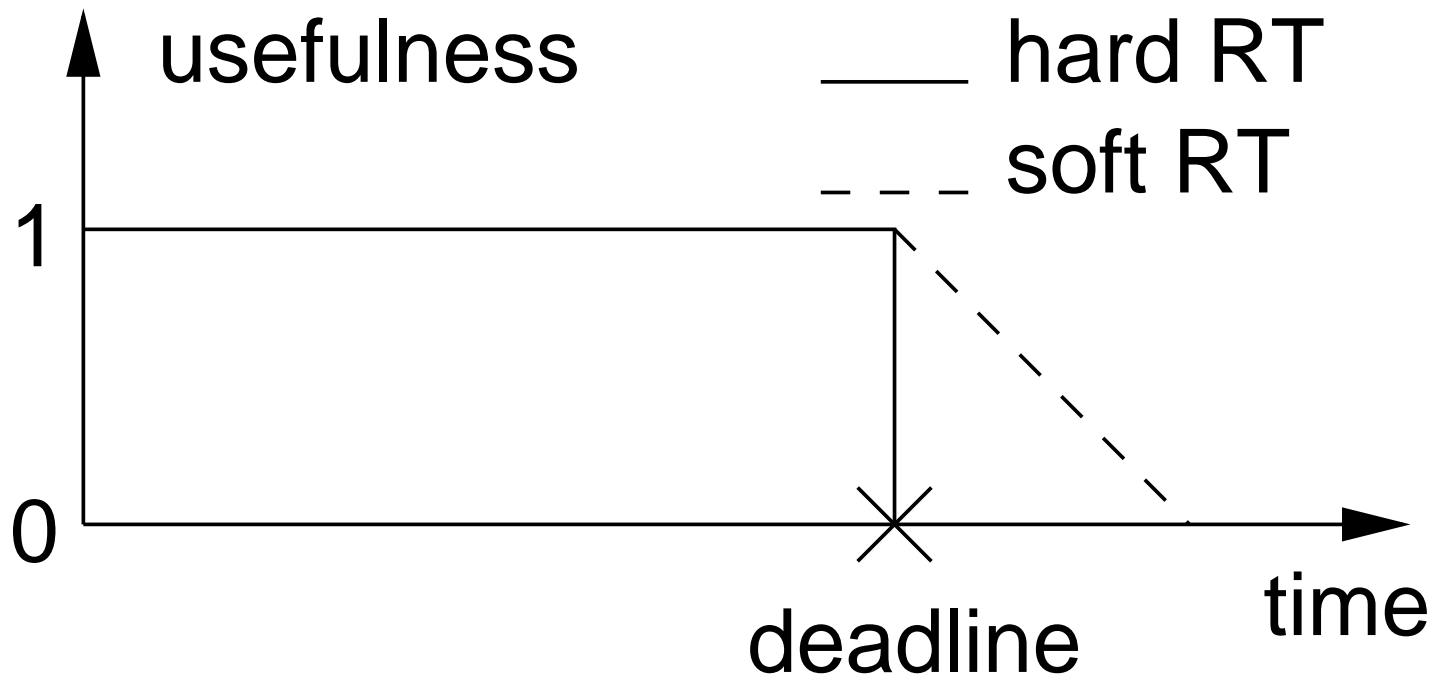
Taxonomy

Dimensions:

- Hard vs. Soft Real-Time
- Time-Triggered vs. Event-Triggered
- Cooperative vs. Preemptive
- Static vs. Dynamic (on/off-line)
- Optimal vs. heuristic


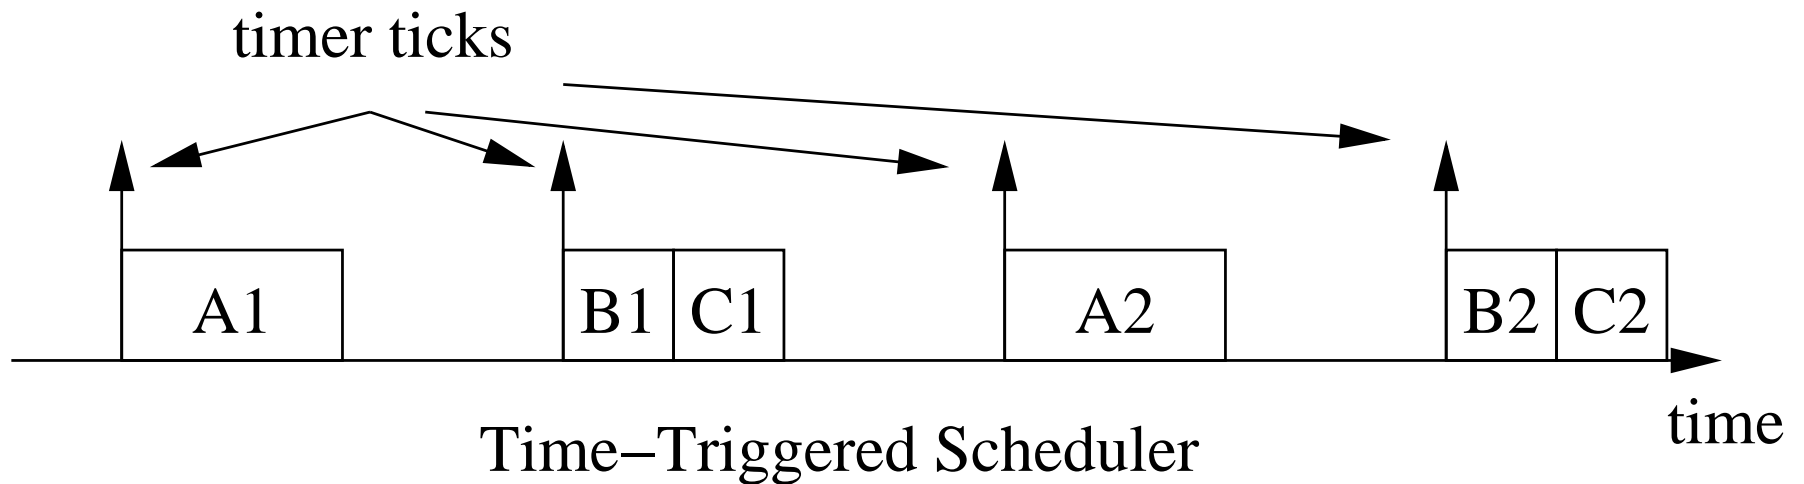
The permutations of the above dimensions give us the taxonomy of schedulers

Real-Time constraints

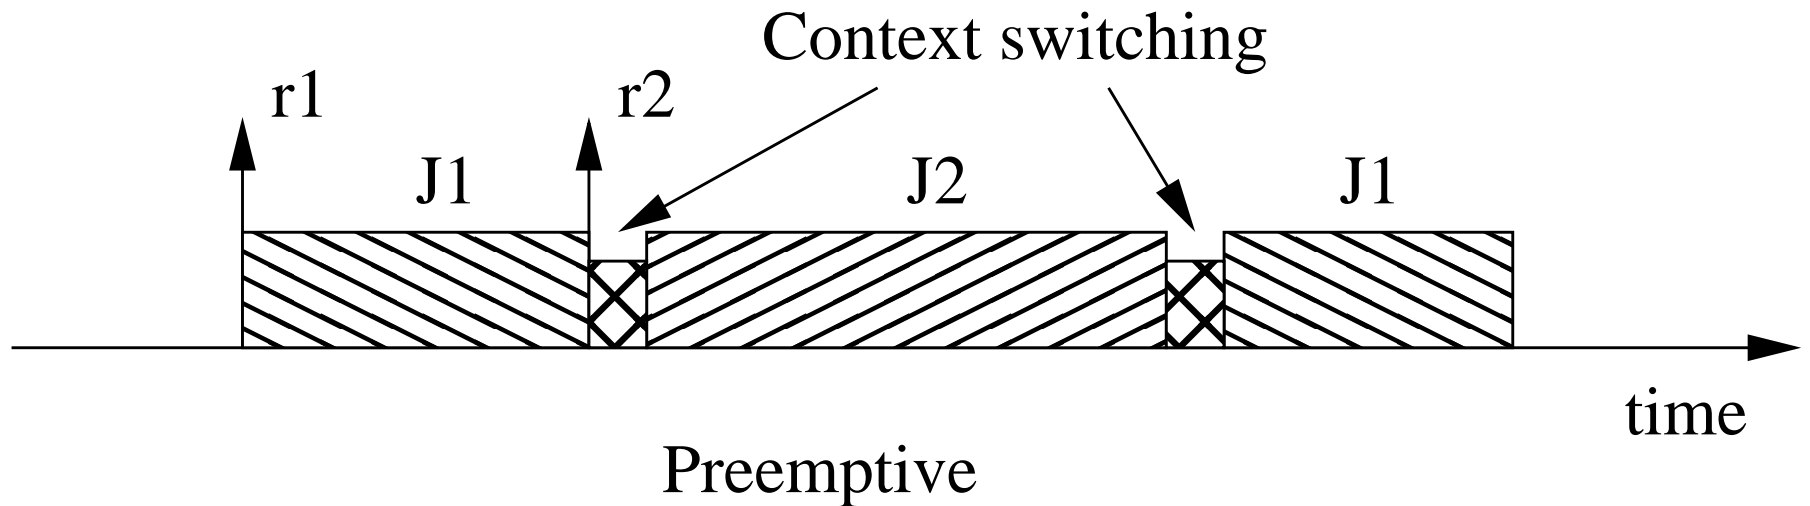



- Value or “usefulness” as a function of time
- It is important to meet the deadline
- High or low performance is not important – “just do it by the deadline”

Time-Triggered vs. Event-Triggered




Cooperative vs. Preemptive




Example: Taxonomy of TTS or ETS

- Preemptive – non-preemptive: preemption means interrupting of one task with the other. Task priorities may be employed.
- Static – dynamic: scheduling decision are based on parameters which can be fixed or not before task activation
- Offline – online: scheduling decisions are made in advance or in the course of system operation
- Optimal – heuristic: An optimal algorithm minimises some given cost function defined for the task set. Feasibility is a simple example of the cost function. Heuristic algorithms use rules, which do not guarantee optimality.

Example: TTS (1,2)




(a) Timeline scheduler



(b) Rate Monotonic scheduler

Example: ETS, EDF, preemptive



Task parameters (1)

- Task J_i or τ_i
- Arrival time a_i or r_i (Release Time, Request) ↑
- Computation Time C_i (exec. without interruption)
- Abs. Deadline d_i ↓
- Rel. Deadline D_i , $D_i = d_i - r_i$ ↓
- Start Time s_i
- Finishing Time f_i

Task parameters (2)

- Response Time R_i , $R_i = f_i - r_i$
- Criticality: hard/soft
- Value v_i
- Lateness L_i , $L_i = f_i - d_i$ (negative is good!)
- Tardiness or Exceeding time E_i , $E_i = \max(0, L_i)$
- Laxity or Slack Time X_i , $X_i = d_i - a_i - C_i$
(the max time a task can be delayed on its activation to complete within its deadline)

Task parameters (3)


- Critical Instant – the time at which the release of the task will produce the largest response time
- Critical Time Zone – the interval between the critical instant and the response time of the corresponding request of the task
- Relative Release Jitter – the max deviation of the start time of two consecutive instances
$$RRJ_i = \max_k |(s_{i,k} - r_{i,k}) - (s_{i,k-1} - r_{i,k-1})|$$
- Absolute Release Jitter – the max deviation of the start time among all instances
$$ARJ_i = \max_k (s_{i,k} - r_{i,k}) - \min_k (s_{i,k} - r_{i,k})$$
- The same for the Execution and Finishing Jitter, Absolute or Relative.

Example: Construct a schedule (1)

Construct an EDF (Earliest Deadline first) schedule. In this table C_i denotes execution time, D_i – relative deadline, T_i – period and τ_i – task name. Use the non-preemptive execution model.

task	C_i	D_i	T_i
τ_1	2	5	6
τ_2	2	4	8
τ_3	4	8	12

Example: Construct a schedule (2)



Example: calculate task parameters

- Absolute Release Jitter *for τ_1 is 2*
- Absolute Finishing Jitter *for τ_2 is 2*
- Absolute Execution Jitter *for τ_3 is 0.*

RM optimality (1)

Read the book! Two “keys” to the proof:


- “Given two periodic tasks τ_1 and τ_2 , with $T_1 < T_2$, if the schedule is feasible by an arbitrary priority assignment, then it is also feasible by RM. That is, RM is optimal.”
[Buttazzo book]

First step: show that a critical instant for any task occurs whenever the task is released simultaneously with all higher-priority tasks.


- Schedulability is checked on critical instants!

Second step: show that if a task set is schedulable by an arbitrary priority assignment, then it is also schedulable by RM.

RM optimality (2)





(a) Different release time S_n and S_{n-1} s — shorter response time R_n




(b) Simultaneous release — longer response time R_n

RM optimality (3)


Consider two cases (the last r_i is in the critical zone of τ_j), try to change the priority, it will not be better



Periodic EDF optimality (1)



(a) Preemptive schedule — optimal



(b) Cooperative schedule — non-optimal and not feasible

Preemption is crucial for optimality!

Periodic EDF optimality (2)

Theorem: A set of periodic tasks is schedulable with EDF iff

$$(1) \quad \sum_{i=1}^n \frac{C_i}{T_i} \leq 1$$

The necessary condition


It is simple – the CPU utilisation cannot exceed 1.

- Create an example!

The sufficient condition

- Proof by contradiction:
 - Assume $U < 1$ and the set of tasks is not schedulable.
- Construct the schedule, find an overrun and find the longest interval $[t_1, t_2]$ of continuous utilisation directly preceding the overrun.
- Observe, that t_1 is the release time for some periodic instance.
- Calculate the total computational time of all tasks in $[t_1, t_2]$.
- Do maths (or just explain) and find the contradiction

The sufficient condition (diag.)



The sufficient condition (maths)

$$(2) \quad C_p(t_1, t_2) = \sum_{r_k \geq t_1, d_k \leq t_2} C_k = \sum_{i=1}^n \left\lfloor \frac{t_2 - t_1}{T_i} \right\rfloor C_i.$$

The key to the proof is the following observation

$$(3) \quad C_p(t_1, t_2) \leq \sum_{i=1}^n \frac{t_2 - t_1}{T_i} C_i = (t_2 - t_1)U.$$

Since a deadline is missed at t_2 , $C_p(t_1, t_2)$ must be greater than the available CPU time $(t_2 - t_1)$. Together with (2) and (3) it gives the following inequality:

$$(4) \quad (t_2 - t_1) < C_p(t_1, t_2) \leq (t_2 - t_1)U.$$

This is not true under the initial assumption of $U < 1$. This concludes the proof.

What is next?

- Presentations
- Please read the book [Buttazzo] – all presentation topics, be prepared to discuss the presentations
- Presentations: 15 min. each, additional time will be booked the next week.
- Please email your presentation to me, use your last name as the name of the file.